

0. Contents

List of Tables	2	7	POSIX	13	
List of Figures	2	7.1	Name	13	
1	Intro to Linux Security	3	7.2	Overview	13
1.1	Do I need to worry	3	7.3	Some Dude on Soflw	14
1.2	The “Hacker” Word	3	8	Unix	16
1.3	Security and Linux	4	9	Linux Kernel vs Os	16
2	Firewalls - The First Line	4	10	References	17
2.1	What is a Firewall	4			
2.2	How a firewall works	5			
2.2.1	Stealth Mode - Discarding Pings	5			
2.2.2	Port Forwarding and Blocking .	5			
2.2.3	Packet Filtering	5			
3	Understanding Services	6			
3.1	Web Server - httpd - Port 80	6			
3.2	Remote Login - telnet - Port 25	6			
3.3	Secure Remote Login - ssh - Port 22 . .	6			
3.4	File Transfer - ftp - Port 21	6			
3.5	Mail Transfer - SMTP - Port 25	6			
4	Linux Firewall - 2nd Line	7			
4.1	The lokkit Command	7			
5	Linux Wireless Security	8			
5.1	Intro to Wireless Security	8			
5.2	What is Encryption ?	8			
6	File System	9			
6.1	Origin of the Term	9			
6.2	Architechture	9			
6.3	Types of File Systems	10			
6.3.1	Flash File Systems	10			
6.3.2	Tape File Systems	10			
6.3.3	Transactional File Systems . . .	11			
6.3.4	Network File Systems	12			
6.3.5	Shared Disk File System	12			
6.3.6	Device File System	12			
6.3.7	Minimal File System	12			

0. List of Tables**0. List of Figures**

1	Firewalls ?	5
---	-----------------------	---

1. Intro to Linux Security

1.1. Do I need to worry

This is certainly a valid question. Your system is one of tens of millions of computers connected to the internet. You aren't a high profile bank that is likely to be targeted by criminals looking for bank account numbers. Should you really worry? After all, how will the hackers possibly find your system amongst all the others?

Let's explore this briefly and see if we need to take steps to protect our system. Like most people I have a DSL internet connection into my home provided by the local telephone company. They have assigned me an Internet Protocol (IP) address that distinguishes me from other users on the internet and supplied me with a DSL modem that is connected to the phone line.

So that everyone in my family can gain access to the internet connection without having to run network cables throughout the house I have a wireless network. This consists of a wireless router/base station that is connected to the DSL modem. The base station is a fairly common low cost device that, like most routers, includes a firewall that provides the first line of defense for my home network. I have configured my wireless network to use the highest level of encryption so the wireless transmissions are as safe as I can make them with current consumer grade technology.

Has anyone found my IP address and tried to get into my network? First we need to talk a little about the way a possible attack might begin – don't worry, we'll cover these topics in greater detail later. Computer systems talk to each through "ports". Specific applications are configured to talk to other systems through specific ports. For example computers might transfer files between each other using something called ftp (File Transfer Protocol). The ftp client and server talk to each other through port 21. The telnet command that allows users to log into one system from another over a network does so over port 25. In fact a Linux system has 65,535 ports that can be used for various forms of communication between different systems on a network. Most of these ports on a Linux system are closed by default – but some are left open simply because closing them renders the system inaccessible to anyone except the person at the keyboard.

It is not surprising to learn, therefore, that the first thing a potential intruder will try to do is see if any useful ports are open. When intruders find my IP address (usually by running a program that tries every IP address known to man until they get a response) they will scan a range of ports to see if any of them are open. So, has anyone tried to find an open port

on my system? The firewall in my wireless router has a log file I can check using a web browser. In an 8 hour period the firewall logged 130 attempts to find an open port to enter my home network. The log file is full of entries that read:

2005/07/15 06:17:45 Connection attempt to base
station from WAN blocked – src:xxx.xxx.xx.x:23648
dst:jnn.nnn.nnn:3306;

Each of these lines represents an attempt to break through the firewall and into a system on my network. Note that the IP addresses have been removed in the above log entry to protect the innocent (as the author of a book on Linux security it would be unwise to publish my IP address) and also, ironically, to protect the guilty (the IP address of the person trying to break into my system from outside).

Now let's take this one step further. I work outside my home from time to time and often need remote access to the Linux server on my home network while on the road. I do this using something called ssh (Secure Shell). The ssh utility is used to remotely log into a computer system from another system. ssh uses port 22. Needless to say I have port 22 open on my firewall and people have found this open port. I checked the logs on my Linux system to find any failed attempts to log in. I found 20 attempts to log in. All these attempts failed because invalid login and password information were entered.

Based on these experiences on what is probably a typical Linux configuration it is safe to say that no matter who you are, as long as you are connected to the internet, either directly via a cable modem or indirectly via a router there is a very good chance you will not escape the attention of those who make it their business, for what ever reason, to try to break into other computer systems.

1.2. The "Hacker" Word

The term "Hacker" is frequently used in the media when describing an individual who breaks into other people's computer systems. This is actually a misuse of a word that at one time did not have the negative connotations it now has. Years ago the word hacker was used to describe a talented computer programmer. Hackers were generally admired for their ability to rapidly write complex and efficient computer programs.

Sadly the term is now used in a derogatory sense to refer to what is essentially a criminal act and this new use for the word is now firmly rooted in popular culture. It is used by the news media, included in book titles and was even adopted by Hollywood the 1996 Angelina Jolie film titled "Hackers". For better or worse the new meaning is here to stay even if those of us who remember the old meaning of the word wish it wasn't so.

106 In this book we will bow to popular culture and use
107 “hacker” in its new context with sincere apologies to
108 those who would prefer that a hacker was still nothing
109 more than a great programmer.

1.3. Security and Linux

110 As Linux users we have some inherent advantages over
111 our fellow Windows users when it comes to security (or
112 lack there of). Hackers, rather like gamblers, use the
113 laws of odds and averages in their endeavors to find
114 vulnerable computer systems to break into. They will
115 typically target the types of systems that have the
116 most security vulnerabilities. They will also mostly
117 focus attention on areas where there are the most op-
118 portunities for unprotected systems – in other words
119 the types are system that are most common on the
120 internet. In both these cases Windows is the predom-
121 inant operating system. In security circles they say
122 Windows has a larger “surface area” to attack both
123 in terms of vulnerabilities and numbers of systems.

124 Linux is both more secure and less common than
125 Windows based systems with the consequence that
126 attacks on Linux systems occur less frequently than
127 on Windows systems. Having said that it would be
128 foolish to be complacent about securing any system
129 regardless of whether it runs Windows, Linux or any
130 other operating system.

131 The purpose of this book is to provide a step by
132 step approach to securing a Linux system from outside
133 attack. It is designed to be used and understood by
134 both new and experienced Linux users.

-\\$-

2. Firewalls - The First Line

135 The most important first step in developing a
136 secure environment is to avoid, wherever pos-
137 sible, having your Linux system being the first
138 line of defense from outside attack. The best way to
139 do this is to ensure that you have a firewall installed
140 between your Linux system (or the network on which
141 it is installed) and the connection to the internet. If,
142 for example, your Linux system is currently connected
143 directly to a cable or DSL modem box then you will
144 need to think seriously about installing a router or
145 wireless base station that includes a firewall feature
146 between the modem and your Linux system. Linux
147 does come with a firewall that can be configured to
148 protect you and we cover this later in the book. It is
149 better, however, not to rely solely on this.

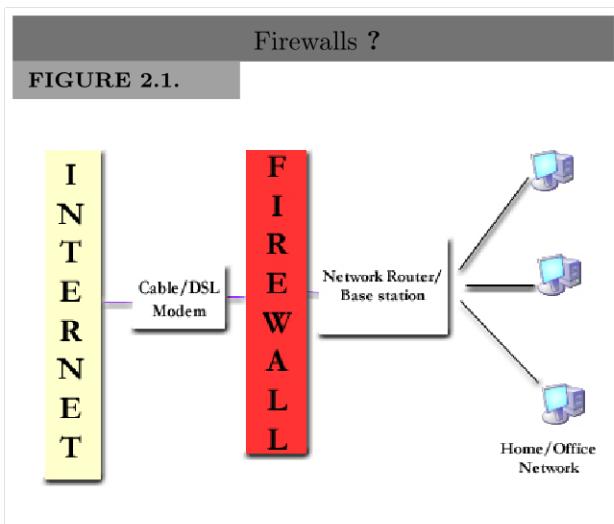
150 A firewall essentially stands between your computer
151 or network on which your computer resides and shields
152 it from the dangers lurking on the internet. It can ei-
153 ther be a software program that runs on a computer
154 system or it can be built into a hardware device such
155 as a wireless base station or router hub. In this chap-
156 ter we are going to look at firewalls as a part of a
157 wired or wireless hub. In later chapters we will look
158 at configuring the firewall software on a Linux system
159 to provide a second layer of defense against attack.

2.1. What is a Firewall

160 It is refreshing in an industry filled with cryptic terms
161 and three letter acronyms (better known as TLAs)
162 to come across a name that is at least somewhat self
163 explanatory.

164 Consider the internet to be inferno of viruses, hack-
165 ers and cyber criminals all looking for systems that
166 they can invade. Rather like a firewall in real life it
167 stops any of this unpleasantness from spreading into
168 your environment. Another good analogy is that of a
169 fortress wall that protects the inhabitants that live in-
170 side. The firewall stops unwanted connections from
171 entering your internal network much like the wall around
172 a fortress prevented the marauding hoards in medieval
173 times. Rather like the gate in the fortress wall the
174 firewall allows only data and connections that meet
175 certain criteria to pass through the wall to the internal
176 network.

177 A typical firewall configuration is shown in Figure
178 2.1. The firewall is positioned between the outside
179 internet connection coming in through the modem
180 and the internal network on which reside a number
181 of Linux and Windows systems. The firewall controls
182 all data traffic and filters out anything that is not
183 permitted to enter the internal network.



2.2. How a firewall works

A typical firewall can perform a number of tasks depending on the complexity of the firewall itself. The basic functions of a firewall are as follows:

2.2.1. Stealth Mode - Discarding Pings

This requires a little explanation. There is a common mechanism in networked environments for finding out if a particular system is up and running and connected to the network. Typically a utility called ping is given the IP address of the remote system. The ping utility sends a data packet to the remote system represented by the IP address and waits for a reply. If it gets a reply then the user knows that the system at that address is available on the network.

Whilst this seems innocuous enough there is actually good reason to configure your firewall to not respond to ping requests. You've probably seen the old war movies (and some new ones too) where the destroyer on the surface of the ocean uses sonar to try to locate a submarine somewhere in the depths below. The sonar sends out pings and waits to see if the sounds bounces back off the hull of the submarine. When the destroyer gets an echo it drops depth charges in an attempt to destroy the submarine. Compare this to your Linux system. The hacker will send out ping packets to every IP address on the planet and attack those that reply. By not responding to the ping packet you have a greater chance of remaining anonymous to the attacker – rather like a stealth submarine that is impervious to sonar.

Don't be fooled by "experts" who try to tell you that ping stands for Packet Internet Groper. This just an attempt by those experts to make something sound more complicated than it is. The author of ping states that he chose that name because of the noise made by sonar.

2.2.2. Port Forwarding and Blocking

Port blocking is the most fundamental level of firewall security and will be used by most home or small business users to protect their systems.

As we mentioned previously computer systems communicate through ports. A firewall can be used to block any ports that you do not want to be open to your systems inside the firewall. For example FTP operates through port 21. If you do not wish anyone on the outside to have ftp access to your systems you will need to configure your firewall to block port 21.

Conversely, Port Forwarding is also a very useful tool to have. Suppose you have three Linux systems on your internal network and want to be able to telnet into one of those systems when you are outside your firewall (perhaps at the local café using the free Wi-Fi connection while you drink your coffee or while in a hotel on a business trip). In this situation you will configure your firewall to forward port 21 connections to the system you want to access from outside. When you connect to your IP address using telnet the firewall will see the packets arriving on port 21 and know that it must forward them to the IP address of the machine you have designated. If you have more than one system on your network it is essential that you set up port forwarding to handle this. After all, without port forwarding how would the router know which internal system you wanted to connect to?

2.2.3. Packet Filtering

Packet filtering is a much more advanced mechanism for providing security and is not available in typical small business or home use router devices.

Data is transmitted over networks and the internet in what are called packets. Each packet contains information about where the data came from and where it is going to (i.e the IP address of the sender and the your IP address). In fact a packet contains a great deal of information about the nature of the data being transmitted and many advanced firewall solutions allow you to filter the data packets coming in through your internet connection to allow or disallow packets depending on what are called filtering rules. For example you might allow a telnet session (which allows you to log into your Linux system from outside) but disallow ftp packets (which allow files to be transferred to and from of your Linux system). You may also choose to block packets arriving from an IP address that you know to be suspicious.

3. Understanding Services

In previous chapters we have touched on some of the services that a Linux system provides and the ports that those services communicate through. In this chapter we will provide an overview of the various communication related services. This information will make it easier to make an informed decision as to whether these are services you want to have running on your Linux system and, therefore, potentially accessible to the outside world.

3.1. Web Server - httpd - Port 80

The httpd service is the Hyper Text Transfer Protocol Deamon. If you plan to host your own web site on your Linux system you will need to activate this service. Without out it your web server will not serve any web pages.

http work through port 80 so you will need to make sure that you have this port open on your Firewall and configured to forward requests to the IP address of the Linux system on your network that is running the web server.

3.2. Remote Login - telnet - Port 25

The telnet service allows users to log into the Linux system from outside. For example you may want to be able to log into your Linux system to perform tasks when you are outside your office or home. You can also use telnet to log into one computer from another on the same network.

The telnet service communicates through port 21. Security experts now advise against the use of telnet these days. Telnet transmits data in plain readable text, which is readily intercepted by hackers leaving vital information (including login and password information) exposed to interception. These days SSH (Secure Shell) is recommended instead.

3.3. Secure Remote Login - ssh - Port 22

Rather like the telnet service the ssh (Secure Shell) service allows users to log into the Linux system from outside. The difference being that ssh uses an encryption mechanism to protect the information being passed over the network thereby preventing others from capturing your login and password information.

The ssh service communicates through port 22.

3.4. File Transfer - ftp - Port 21

FTP is short for File Transfer Protocol and is the protocol for exchanging files over the Internet. FTP is most commonly used to download a file from a server using the Internet or to upload a file to a server. FTP uses port 21 so if you think you or others will need to transfer files to or from your Linux system make sure port 21 is configured correctly on your Firewall.

The vsftp (very secure ftp) server is recommended since it is more secure than the standard ftp server. It is also considered to be smaller and faster.

3.5. Mail Transfer - SMTP - Port 25

SMTP is short for Simple Mail Transfer Protocol and is a protocol for sending e-mail messages between servers. Most e-mail systems that send mail over the Internet use SMTP to send messages from one server to another. The messages can then be retrieved with an e-mail client such as Evolution, KMail, or Balsa.

-§-

4. Linux Firewall - 2nd Line

world in some way (perhaps through a router or firewall to a broadband connection). The firewall feature allows you to disable the firewall settings for any connections coming in from the device connected to the trusted or secure network while applying the firewall rules to device connected ot the untrusted network.

-§-

364
365
366
367
368
369

319 In previous chapters we have covered the firewall
320 located in the router or cable modem and viewed
321 this as the first line of defense in protecting your
322 Linux system from outside attack. In this chapter we
323 will be looking at the second line of defense – the
324 firewall on your Linux system.

325 During the installation of your Linux system you
326 will have been asked a number of questions about the
327 security settings you wanted to select. At the time you
328 may not have understood what these settings meant
329 or you may not recall which settings you chose. In this
330 Chapter we will explore how to configure the security
331 settings of your Linux system.

4.1. The lokkit Command

332 The lokkit command can be run at any time to change
333 the security settings of Firewall installed on your sys-
334 tem. To run this command you must first login as root
335 or use the “su” command. If you are already super
336 user on your Linux system start the lokkit command
337 as follows:

338 /usr/sbin/lokkit

339 or to use the su command from a non-super user
340 account as follows:

341 su {c “/usr/sbin/lokkit”

342 The lokkit command allows you to either enable or
343 disable the Firewall. The first step if it is not already
344 enabled is to enable it. Use the “Tab” key to move
345 around and the “Space” key to select the “Enabled”
346 option.

347 The second step is configure the Firewall. Use the
348 Tab key to move the “Configure” button and press
349 the “Space” key.

350 On the configuration screen simply select the ser-
351 vice types that you want to support. Based on your
352 selections lokkit will configure the Firewall to allow
353 access to the appropriate ports. The services listed
354 are HTTP, FTP, SSH, Telnet and Mail (SMTP). You
355 can also specify other ports you wish to open on the
356 Firewall in the “other ports” section.

357 The lokkit command also provides the option of
358 specifying trusted devices on the “Configure” screen.
359 In summary, it is possible to have more than one net-
360 work device installed on a Linux system. In this sce-
361 nario it might be that one device is connected to a
362 trusted and secure network while the other is con-
363 nected to a network that is connected to the outside

364
365
366
367
368
369

5. Linux Wireless Security

The saying goes that a chain is only as strong as its weakest link and network security is no exception to this rule. In previous chapters we've looked a number of places where you can improve the level of security of your Linux system. We will now look at another area of Linux security that, if not implemented correctly, will make all the other security precautions ineffective. This is the area of wireless security.

Home and business networks are now increasingly going wireless. This has many benefits primarily in terms of convenience. Wireless networks mean that network cable doesn't have to run wherever a computer needs to be installed and you can get network and internet access anywhere as long as you are within the range of wireless base station. This, for example, gives users freedom to use their laptops in places in a building where there is no network connection available.

There is one big draw back to wireless networks in that they provide another point of security vulnerability in the network. No matter how secure you have made you have made your firewall and your Linux system if you have not also secured your wireless network a hacker can very easily eaves drop and monitor all the traffic on your network to get information such as system passwords and bank account login and password information. Even worse, if your wireless network is wide open intruders can connect to your network and potentially access all of your systems.

Carefully configuring your firewall and Linux system whilst leaving your wireless network unprotected is akin to locking the front door of your house but leaving the windows open.

Fortunately securing a wireless network is straightforward and can be implemented with a few simple steps. In this chapter we will explore the world of wireless security and show you how to ensure your wireless network is as secure as it can be made with current wireless technology.

5.1. Intro to Wireless Security

Wireless data differs from data traveling through a wired network in that the data is broadcast using radio waves. These radio transmissions pass through walls and floors and ceilings into the apartments above or below, the street outside or the house or office building next door. While data traveling through an ethernet cable is almost impossible to intercept the data from a WiFi network can potentially be picked up by anyone with a wireless network card within the range of the wireless network.

In the wired world we rely on firewalls to protect networks and systems from intrusion. The wireless network is typically located behind the firewall and attack comes not from a hacker attempting to break in through your internet connection but from a person in the building or room next door or the opportunistic hacker who drives the streets at night with a laptop looking for unprotected wireless networks.

Wireless networks are protected from attack by using encryption. This ensures that the data passing between the computers on the network and the wireless base station/router can only be understood by other computers that know what key was used to encrypt the data. It is very unlikely that a hacker will be able to find out what your encryption key is. In fact breaking into encrypted wireless networks is so difficult and time consuming that the hacker will simply take the path of least resistance and move on to one of the many unprotected wireless networks rather than try to break into yours.

There is no practical way to prevent these radio waves carrying our data from spreading outside our buildings (short of encasing them in lead) so we have to accept that the data is going to be visible to others. Rather than preventing the data from being seen by others, therefore, we instead rely on encryption to make the data unintelligable to the hacker. Whilst anyone in range of our wireless network can see the data they cannot read it without the correct encryption key.

5.2. What is Encryption ?

Encryption essentially involves taking data and subjecting it to mathematical algorithms that include a key making it unreadable to anyone else who does not know what that key is. The encrypted form of the data is known as ciphertext. Wireless networks use what is known as symmetrical encryption whereby the same key is used at both ends of the network connection. For example, the encryption key is used as part of the mathematical equation on the sending system to encrypt the data. The receiving system then uses the same key to decrypt the data when it receives it. This key is specified by you when you configure the encryption for your wireless network and should be known only to you. The chances of a hacker guessing your encryption key are very remote and while it is possible to break the encryption code with enough time and computing power it is unlikely this kind of effort will be expended on your network. You can specify different lengths of key for the encryption process - the longer the key the stronger the encryption and the more secure the network.

WiFi wireless networks use a security standard known as Wired Equivalent Privacy (WEP). The aim of WEP is to provide a level of security in a wireless network environment that is equivalent to the security of a wired network. In practice it falls short of this

476 goal but for most purposes it provides an adequate
477 level of protection.

478 Wireless encryption can be configured as either 64-
479 bit or 128-bit. This refers to the length of the key that
480 is used in the encryption algorithm and these relate
481 directly to the strength of the encryption (128-bit en-
482 cryption being stronger than 64-bit encryption). Using
483 stronger encryption can impact the performance of the
484 network because more time has to be spent encrypting
485 and decrypting the data at each end of the communi-
486 cation. In practice it is unlikely the typical user would
487 notice a significant difference and the strongest encryp-
488 tion (128-bit) is always recommended.

489 The encryption key are specified in hexadecimal.
490 Unlike decimal which uses a number base of 10 (i.e.
491 digits between 0 - 9) hexadecimal uses a base of 16
492 (i.e digits between 0 - 9 and A - F). 64-bit encryption
493 requires that you provide a 10 digit key whilst 128-bit
494 encryption requires that you provide a 26 digit key.

-§-

6. File System

In computing, a file system or filesystem (often ab-
495 breviated to fs), controls how data is stored and
496 retrieved. Without a file system, data placed in a
497 storage medium would be one large body of data with
498 no way to tell where one piece of data stops and the
499 next begins. By separating the data into pieces and
500 giving each piece a name, the data is easily isolated
501 and identified. Taking its name from the way paper-
502 based data management system is named, each group
503 of data is called a “file”. The structure and logic rules
504 used to manage the groups of data and their names is
505 called a “file system”.

There are many different kinds of file systems. Each
507 one has different structure and logic, properties of
508 speed, flexibility, security, size and more. Some file
509 systems have been designed to be used for specific
510 applications. For example, the ISO 9660 file system is
511 designed specifically for optical discs.

File systems can be used on numerous different
513 types of storage devices that use different kinds of
514 media. As of 2019, hard disk drives have been key
515 storage devices and are projected to remain so for the
516 foreseeable future. Other kinds of media that are used
517 include SSDs, magnetic tapes, and optical discs. In
518 some cases, such as with tmpfs, the computer’s main
519 memory (random-access memory, RAM) is used to cre-
520 ate a temporary file system for short-term use.

Some file systems are used on local data storage de-
522 vices; others provide file access via a network protocol
523 (for example, NFS, SMB, or 9P clients). Some file sys-
524 tems are “virtual”, meaning that the supplied “files
525 ” (called virtual files) are computed on request (such
526 as procfs and sysfs) or are merely a mapping into a
527 different file system used as a backing store. The file
528 system manages access to both the content of files and
529 the metadata about those files. It is responsible for
530 arranging storage space; reliability, efficiency, and tun-
531 ing with regard to the physical storage medium are
532 important design considerations.

6.1. Origin of the Term

Before the advent of computers the term file system
534 was used to describe a method of storing and retrieving
535 paper documents. By 1961 the term was being applied
536 to computerized filing alongside the original meaning.
537 By 1964 it was in general use.

6.2. Architecture

A file system consists of two or three layers. Sometimes
539 the layers are explicitly separated, and sometimes the
540 functions are combined.

The logical file system is responsible for interaction with the user application. It provides the application program interface (API) for file operations — OPEN, CLOSE, READ, etc., and passes the requested operation to the layer below it for processing. The logical file system “ manage[s] open file table entries and per-process file descriptors. ” This layer provides “ file access, directory operations, [and] security and protection. ”

The second optional layer is the virtual file system. “ This interface allows support for multiple concurrent instances of physical file systems, each of which is called a file system implementation. ”[8]

The third layer is the physical file system. This layer is concerned with the physical operation of the storage device (e.g. disk). It processes physical blocks being read or written. It handles buffering and memory management and is responsible for the physical placement of blocks in specific locations on the storage medium. The physical file system interacts with the device drivers or with the channel to drive the storage device.[7]

6.3. Types of File Systems

File system types can be classified into disk/tape file systems, network file systems and special-purpose file systems. Disk file systems

A disk file system takes advantages of the ability of disk storage media to randomly address data in a short amount of time. Additional considerations include the speed of accessing data following that initially requested and the anticipation that the following data may also be requested. This permits multiple users (or processes) access to various data on the disk without regard to the sequential location of the data. Examples include FAT (FAT12, FAT16, FAT32), ex-FAT, NTFS, HFS and HFS+, HPFS, APFS, UFS, ext2, ext3, ext4, XFS, btrfs, ISO 9660, Files-11, Veritas File System, VMFS, ZFS, ReiserFS and UDF. Some disk file systems are journaling file systems or versioning file systems. Optical discs

ISO 9660 and Universal Disk Format (UDF) are two common formats that target Compact Discs, DVDs and Blu-ray discs. Mount Rainier is an extension to UDF supported since 2.6 series of the Linux kernel and since Windows Vista that facilitates rewriting to DVDs.

6.3.1. Flash File Systems

Flash file systems

Main article: Flash file system

A flash file system considers the special abilities, performance and restrictions of flash memory devices. Frequently a disk file system can use a flash memory device as the underlying storage media but it is much

better to use a file system specifically designed for a flash device.

6.3.2. Tape File Systems

Tape file systems

A tape file system is a file system and tape format designed to store files on tape in a self-describing form[clarification needed]. Magnetic tapes are sequential storage media with significantly longer random data access times than disks, posing challenges to the creation and efficient management of a general-purpose file system.

In a disk file system there is typically a master file directory, and a map of used and free data regions. Any file additions, changes, or removals require updating the directory and the used/free maps. Random access to data regions is measured in milliseconds so this system works well for disks.

Tape requires linear motion to wind and unwind potentially very long reels of media. This tape motion may take several seconds to several minutes to move the read/write head from one end of the tape to the other.

Consequently, a master file directory and usage map can be extremely slow and inefficient with tape. Writing typically involves reading the block usage map to find free blocks for writing, updating the usage map and directory to add the data, and then advancing the tape to write the data in the correct spot. Each additional file write requires updating the map and directory and writing the data, which may take several seconds to occur for each file.

Tape file systems instead typically allow for the file directory to be spread across the tape intermixed with the data, referred to as streaming, so that time-consuming and repeated tape motions are not required to write new data.

However, a side effect of this design is that reading the file directory of a tape usually requires scanning the entire tape to read all the scattered directory entries. Most data archiving software that works with tape storage will store a local copy of the tape catalog on a disk file system, so that adding files to a tape can be done quickly without having to rescan the tape media. The local tape catalog copy is usually discarded if not used for a specified period of time, at which point the tape must be re-scanned if it is to be used in the future.

IBM has developed a file system for tape called the Linear Tape File System. The IBM implementation of this file system has been released as the open-source IBM Linear Tape File System — Single Drive Edition (LTFS-SDE) product. The Linear Tape File System uses a separate partition on the tape to record the index meta-data, thereby avoiding the problems associated with scattering directory entries across the

647 entire tape.

648 Tape formatting

649 Writing data to a tape, erasing, or formatting a tape
650 is often a significantly time-consuming process and can
651 take several hours on large tapes.[a] With many data
652 tape technologies it is not necessary to format the
653 tape before over-writing new data to the tape. This is
654 due to the inherently destructive nature of overwriting
655 data on sequential media.

656 Because of the time it can take to format a tape,
657 typically tapes are pre-formatted so that the tape user
658 does not need to spend time preparing each new tape
659 for use. All that is usually necessary is to write an
660 identifying media label to the tape before use, and
661 even this can be automatically written by software
662 when a new tape is used for the first time. Database
663 file systems

664 Another concept for file management is the idea of
665 a database-based file system. Instead of, or in addition
666 to, hierarchical structured management, files are iden-
667 tified by their characteristics, like type of file, topic,
668 author, or similar rich metadata.[12]

669 IBM DB2 for i [13] (formerly known as DB2/400
670 and DB2 for i5/OS) is a database file system as part
671 of the object based IBM i [14] operating system (for-
672 merly known as OS/400 and i5/OS), incorporating a
673 single level store and running on IBM Power Systems
674 (formerly known as AS/400 and iSeries), designed by
675 Frank G. Soltis IBM's former chief scientist for IBM
676 i. Around 1978 to 1988 Frank G. Soltis and his team
677 at IBM Rochester have successfully designed and ap-
678 plied technologies like the database file system where
679 others like Microsoft later failed to accomplish.[15]
680 These technologies are informally known as Fortress
681 Rochester[citation needed] and were in few basic as-
682 pects extended from early Mainframe technologies but
683 in many ways more advanced from a technological per-
684 spective[citation needed].

685 Some other projects that aren't "pure" database
686 file systems but that use some aspects of a database
687 file system:

688 Many Web content management systems use a rela-
689 tional DBMS to store and retrieve files. For example,
690 XHTML files are stored as XML or text fields, while
691 image files are stored as blob fields; SQL SELECT
692 (with optional XPath) statements retrieve the files,
693 and allow the use of a sophisticated logic and more rich
694 information associations than "usual file systems".
695 Many CMSs also have the option of storing only meta-
696 data within the database, with the standard filesystem
697 used to store the content of files. Very large file sys-
698 tems, embodied by applications like Apache Hadoop
699 and Google File System, use some database file system
700 concepts.

6.3.3. Transactional File Systems

701 Transactional file systems

702 Some programs need to either make multiple file
703 system changes, or, if one or more of the changes fail
704 for any reason, make none of the changes. For example,
705 a program which is installing or updating software may
706 write executables, libraries, and/or configuration files.
707 If some of the writing fails and the software is left
708 partially installed or updated, the software may be
709 broken or unusable. An incomplete update of a key
710 system utility, such as the command shell, may leave
711 the entire system in an unusable state.

712 Transaction processing introduces the atomicity
713 guarantee, ensuring that operations inside of a trans-
714 action are either all committed or the transaction can
715 be aborted and the system discards all of its partial
716 results. This means that if there is a crash or power
717 failure, after recovery, the stored state will be consis-
718 tent. Either the software will be completely installed
719 or the failed installation will be completely rolled back,
720 but an unusable partial install will not be left on the
721 system. Transactions also provide the isolation guar-
722 antee[clarification needed], meaning that operations
723 within a transaction are hidden from other threads on
724 the system until the transaction commits, and that
725 interfering operations on the system will be properly
726 serialized with the transaction.

727 Windows, beginning with Vista, added transaction
728 support to NTFS, in a feature called Transactional
729 NTFS, but its use is now discouraged.[16] There are
730 a number of research prototypes of transactional file
731 systems for UNIX systems, including the Valor file
732 system,[17] Amino,[18] LFS,[19] and a transactional
733 ext3 file system on the TxOS kernel,[20] as well as
734 transactional file systems targeting embedded systems,
735 such as TFFS.[21]

736 Ensuring consistency across multiple file system op-
737 erations is difficult, if not impossible, without file sys-
738 tem transactions. File locking can be used as a con-
739 currency control mechanism for individual files, but it
740 typically does not protect the directory structure or
741 file metadata. For instance, file locking cannot prevent
742 TOCTTOU race conditions on symbolic links. File
743 locking also cannot automatically roll back a failed
744 operation, such as a software upgrade; this requires
745 atomicity.

746 Journaling file systems is one technique used to in-
747 troduce transaction-level consistency to file system
748 structures. Journal transactions are not exposed to
749 programs as part of the OS API they are only used
750 internally to ensure consistency at the granularity of
751 a single system call.

752 Data backup systems typically do not provide sup-
753 port for direct backup of data stored in a transactional
754 manner, which makes the recovery of reliable and con-
755 sistent data sets difficult. Most backup software simply
756 notes what files have changed since a certain time, re-
757 gardless of the transactional state shared across multi-

758 ple files in the overall dataset. As a workaround, some
 759 database systems simply produce an archived state file
 760 containing all data up to that point, and the backup
 761 software only backs that up and does not interact di-
 762 rectly with the active transactional databases at all.
 763 Recovery requires separate recreation of the database
 764 from the state file after the file has been restored by
 765 the backup software.

6.3.4. Network File Systems

766 Main article: Distributed file system

767 A network file system is a file system that acts as
 768 a client for a remote file access protocol, providing
 769 access to files on a server. Programs using local
 770 interfaces can transparently create, manage and access
 771 hierarchical directories and files in remote network-
 772 connected computers. Examples of network file sys-
 773 tems include clients for the NFS, AFS, SMB protocols,
 774 and file-system-like clients for FTP and WebDAV.

6.3.5. Shared Disk File System

775 Shared disk file systems

776 Main article: Shared disk file system

777 A shared disk file system is one in which a number
 778 of machines (usually servers) all have access to the
 779 same external disk subsystem (usually a SAN). The
 780 file system arbitrates access to that subsystem, pre-
 781 venting write collisions. Examples include GFS2 from
 782 Red Hat, GPFS from IBM, SFS from DataPlow, CXFS
 783 from SGI and StorNext from Quantum Corporation.
 784 Special file systems

785 A special file system presents non-file elements of
 786 an operating system as files so they can be acted on
 787 using file system APIs. This is most commonly done
 788 in Unix-like operating systems, but devices are given
 789 file names in some non-Unix-like operating systems as
 790 well.

6.3.6. Device File System

791 Device file systems

792 A device file system represents I/O devices and
 793 pseudo-devices as files, called device files. Examples in
 794 Unix-like systems include devfs and, in Linux 2.6 sys-
 795 tems, udev. In non-Unix-like systems, such as TOPS-
 796 10 and other operating systems influenced by it, where
 797 the full filename or pathname of a file can include a
 798 device prefix, devices other than those containing file
 799 systems are referred to by a device prefix specifying
 800 the device, without anything following it. Other spe-
 801 cial file systems

802 In the Linux kernel, configfs and sysfs provide files
 803 that can be used to query the kernel for information
 804 and configure entities in the kernel. procfs maps pro-
 805 cesses and, on Linux, other operating system struc-

806 tures into a filesystem.

6.3.7. Minimal File System

807 Minimal file system / audio-cassette storage

808 In the 1970s disk and digital tape devices were too
 809 expensive for some early microcomputer users. An in-
 810 inexpensive basic data storage system was devised that
 811 used common audio cassette tape.

812 When the system needed to write data, the user
 813 was notified to press " RECORD " on the cassette
 814 recorder, then press " RETURN " on the keyboard
 815 to notify the system that the cassette recorder was
 816 recording. The system wrote a sound to provide time
 817 synchronization, then modulated sounds that encoded
 818 a prefix, the data, a checksum and a suffix. When the
 819 system needed to read data, the user was instructed to
 820 press " PLAY " on the cassette recorder. The system
 821 would listen to the sounds on the tape waiting until
 822 a burst of sound could be recognized as the synchro-
 823 nization. The system would then interpret subsequent
 824 sounds as data. When the data read was complete, the
 825 system would notify the user to press " STOP " on
 826 the cassette recorder. It was primitive, but it worked
 827 (a lot of the time). Data was stored sequentially, usu-
 828 ally in an unnamed format, although some systems
 829 (such as the Commodore PET series of computers)
 830 did allow the files to be named. Multiple sets of data
 831 could be written and located by fast-forwarding the
 832 tape and observing at the tape counter to find the
 833 approximate start of the next data region on the tape.
 834 The user might have to listen to the sounds to find
 835 the right spot to begin playing the next data region.
 836 Some implementations even included audible sounds
 837 interspersed with the data. Flat file systems

838 Not to be confused with Flat file database.

839 In a flat file system, there are no subdirectories;
 840 directory entries for all files are stored in a single di-
 841 rectory.

842 When floppy disk media was first available this type
 843 of file system was adequate due to the relatively small
 844 amount of data space available. CP/M machines fea-
 845 tured a flat file system, where files could be assigned
 846 to one of 16 user areas and generic file operations nar-
 847 rowed to work on one instead of defaulting to work on
 848 all of them. These user areas were no more than spec-
 849 ial attributes associated with the files; that is, it was
 850 not necessary to define specific quota for each of these
 851 areas and files could be added to groups for as long as
 852 there was still free storage space on the disk. The early
 853 Apple Macintosh also featured a flat file system, the
 854 Macintosh File System. It was unusual in that the file
 855 management program (Macintosh Finder) created the
 856 illusion of a partially hierarchical filing system on top
 857 of EMFS. This structure required every file to have a
 858 unique name, even if it appeared to be in a separate
 859 folder. IBM DOS/360 and OS/360 store entries for
 860 all files on a disk pack (volume) in a directory on the

861 pack called a Volume Table of Contents (VTOC).

862 While simple, flat file systems become awkward as
863 the number of files grows and makes it difficult to
864 organize data into related groups of files.

865 A recent addition to the flat file system family is
866 Amazon's S3, a remote storage service, which is in-
867 tentionally simplistic to allow users the ability to cus-
868 tomize how their data is stored. The only constructs
869 are buckets (imagine a disk drive of unlimited size) and
870 objects (similar, but not identical to the standard con-
871 cept of a file). Advanced file management is allowed by
872 being able to use nearly any character (including '/')
873 in the object's name, and the ability to select subsets
874 of the bucket's content based on identical prefixes.

-§-

7. POSIX

875 **T**he Portable Operating System Interface
876 (POSIX) is a family of standards specified by
877 the IEEE Computer Society for maintaining
878 compatibility between operating systems. POSIX de-
879 fines the application programming interface (API),
880 along with command line shells and utility interfaces,
881 for software compatibility with variants of Unix and
882 other operating systems.

7.1. Name

883 Originally, the name POSIX referred to IEEE Std
884 1003.1-1988, released in 1988. The family of POSIX
885 standards is formally designated as IEEE 1003 and
886 the international standard name is ISO/IEC 9945.

887 The standards emerged from a project that be-
888 gan circa 1985. Richard Stallman suggested the name
889 POSIX to the IEEE instead of former IEEE-IX. The
890 committee found it more easily pronounceable and
891 memorable, and thus adopted it.

7.2. Overview

892 Unix was selected as the basis for a standard system
893 interface partly because it was "manufacturer-neutral"
894 . However, several major versions of Unix existed—so
895 there was a need to develop a common denominator
896 system. The POSIX specifications for Unix-like oper-
897 ating systems originally consisted of a single document
898 for the core programming interface, but eventually
899 grew to 19 separate documents (POSIX.1, POSIX.2,
900 etc). The standardized user command line and script-
901 ing interface were based on the UNIX System V shell.
902 Many user-level programs, services, and utilities (in-
903 cluding awk, echo, ed) were also standardized, along
904 with required program-level services (including basic
905 I/O: file, terminal, and network). POSIX also defines a
906 standard threading library API which is supported by
907 most modern operating systems. In 2008, most parts
908 of POSIX were combined into a single standard (IEEE
909 Std 1003.1-2008, also known as POSIX.1-2008).

910 As of 2014, POSIX documentation is divided into
911 two parts:

912 **POSIX.1, 2013 Edition:** POSIX Base Definitions,
913 System Interfaces, and Commands and Utilities (which
914 include POSIX.1, extensions for POSIX.1, Real-time
915 Services, Threads Interface, Real-time Extensions, Se-
916 curity Interface, Network File Access and Network
917 Process-to-Process Communications, User Portability
918 Extensions, Corrections and Extensions, Protection
919 and Control Utilities and Batch System Utilities. This
920 is POSIX 1003.1-2008 with Technical Corrigendum 1.)
921 **POSIX Conformance Testing:** A test suite for POSIX

922 accompanies the standard: VSX-PCTS or the VSX
 923 POSIX Conformance Test Suite.

924 The development of the POSIX standard takes place
 925 in the Austin Group (a joint working group among the
 926 IEEE, The Open Group, and the ISO/IEC JTC 1).

7.3. Some Dude on Softw

927 The most important things POSIX 7 defines

C API

929 Greatly extends ANSI C with things like:

more file operations :

mkdir
dirname
symlink
readlink
link (hardlinks)
poll()
stat
sync
nftw()

process and threads :

fork
exec
wait
pipe
semaphors sem_*

shared memory (shm_*)
kill

scheduling parameters (nice_*|priched_*)

sleep
mkfifo
setpgid()
socket()
mmap
mlock
mprotect
madvise
brk()

networking :
memory management :

mmap
mlock
mprotect
madvise
brk()

utilities : regular expressions

930 Those APIs also determine underlying system
 931 concepts on which they depend, e.g. fork requires a
 932 concept of a process.

933 Many Linux system calls exist to implement a
 934 specific POSIX C API function and make Linux compliant,
 935 e.g. sys_write, sys_read, ... Many of those syscalls
 936 also have Linux-specific extensions however.

937 Major Linux desktop implementation: glibc, which
 938 in many cases just provides a shallow wrapper to sys-
 939 tem calls.

940 CLI utilities

941 E.g.: cd, ls, echo, ...

942 Many utilities are direct shell front ends for a cor-
 943 responding C API function, e.g. mkdir.

944 Major Linux desktop implementation: GNU Core-
 945 utils for the small ones, separate GNU projects for
 946 the big ones: sed, grep, awk, ... Some CLI utilities are

implemented by Bash as built-ins.

Shell language

E.g., a=b; echo "\$a"

Major Linux desktop implementation: GNU Bash.

Environment variables

E.g.: HOME, PATH.

PATH search semantics are specified, including how slashes prevent PATH search.

Program exit status

ANSI C says 0 or EXIT_SUCCESS for success,
EXIT_FAILURE for failure, and leaves the rest im-
plementation defined.

POSIX adds:

126: command found but not executable.

127: command not found.

& 128: terminated by a signal.

But POSIX does not seem to specify
the 128 + SIGNAL_ID rule used by Bash:
<https://unix.stackexchange.com/questions/99112/default-exit-code-when-process-is-terminated>

* Regular expression

There are two types: BRE (Basic) and ERE (Ex-
tended). Basic is deprecated and only kept to not break
APIs.

Those are implemented by C API functions, and
used throughout CLI utilities, e.g. grep accepts BREs
by default, and EREs with -E.

E.g.: echo 'a.1' — grep -E 'a.[[:digit:]]'

Major Linux implementation: glibc implements the
functions under regex.h which programs like grep can
use as backend.

Directory struture (reg)

E.g.: /dev/null, /tmp

The Linux FHS greatly extends POSIX.

Filenames / is the path separator portable filenames
use at most max 14 chars and 256 for the full path
can only contain: a-zA-Z0-9_-

See also: what is posix compliance for filesystem?

Command line utility API conventions

Not mandatory, used by POSIX, but almost nowhere
else, notably not in GNU. But true, it is too restrictive,
e.g. single letter flags only (e.g. -a), no double hyphen
long versions (e.g. --all).

A few widely used conventions: - means stdin where
a file is expected – terminates flags, e.g. ls -l to list
a directory named -l

See also: Are there standards for Linux command
line switches and arguments?

995 “POSIX ACLs” (Access Control Lists), e.g. as used POSIX-compatible?
996 as backend for setfacl.

997 This was withdrawn but it was implemented in sev-
998 eral OSes, including in Linux with setxattr.

999 Who conforms to POSIX?

1000 Many systems follow POSIX closely, but few are
1001 actually certified by the Open Group which maintains
1002 the standard. Notable certified ones include:

1003 OS X (Apple) X stands for both 10 and UNIX. Was
1004 the first Apple POSIX system, released circa 2001. See
1005 also: Is OSX a POSIX OS? AIX (IBM) HP-UX (HP)
1006 Solaris (Oracle)

1007 Most Linux distros are very compliant, but not cer-
1008 tified because they don't want to pay the compliance
1009 check. Inspur's K-UX and Huawei's EulerOS are two
1010 certified examples.

1011 The official list of certified systems be found at:
1012 <https://www.opengroup.org/openbrand/register/> and
1013 also at the wiki page.

1014 Windows

1015 Windows implemented POSIX on some of its pro-
1016 fessional distributions.

1017 Since it was an optional feature, programmers could
1018 not rely on it for most end user applications.

1019 Support was deprecated in Windows 8:

1020 Where does Microsoft Windows' 7
1021 POSIX implementation currently stand?
1022 [https://superuser.com/questions/495360/does-](https://superuser.com/questions/495360/does-windows-8-still-implement-posix)
1023 windows-8-still-implement-posix Feature request:
1024 [https://windows.uservoice.com/forums/265757-](https://windows.uservoice.com/forums/265757-windows-feature-suggestions/suggestions/6573649-full-posix-support)
1025 windows-feature-suggestions/suggestions/6573649-
1026 full-posix-support

1027 In 2016 a new official Linux-like API
1028 called *Windows Subsystem for Linux* was an-
1029 nounced. It includes Linux system calls, ELF
1030 running, parts of the /proc filesystem, Bash,
1031 GCC, (TODO likely glibc?), apt-get and more:
1032 <https://channel9.msdn.com/Events/Build/2016/P488>
1033 so I believe that it will allow Windows to run much, if
1034 not all, of POSIX. However, it is focused on developers
1035 / deployment instead of end users. In particular, there
1036 were no plans to allow access to the Windows GUI.

1037 Historical overview of the official Microsoft POSIX
1038 compatibility: [http://brianreiter.org/2010/08/24/the-](http://brianreiter.org/2010/08/24/the-sad-history-of-the-microsoft-posix-subsystem/)
1039 [sad-history-of-the-microsoft-posix-subsystem/](http://brianreiter.org/2010/08/24/the-sad-history-of-the-microsoft-posix-subsystem/)

1040 Cygwin is a well known GPL third-party project
1041 for that *provides substantial POSIX API functionality*
1042 for Windows, but requires that you *rebuild your applica-*
1043 *tion from source if you want it to run on Windows.*
1044 MSYS2 is a related project that seems to add more
1045 functionality on top of Cygwin.

1046 Android

1047 Android has its own C library (Bionic) which does
1048 not fully support POSIX as of Android O: Is Android

-§-

8. Unix

1050 Unix was made by two guys Ken Thompson (inven-
1051 tor of utf-8) and Dennis Ritchie (inventor or C). So
1052 these two behemoths were working on a computing sys-
1053 tem called multex : Multiplexed information computer
1054 service.

1055 The objective of this was to have multiple programs
1056 running at the same time, hence multiplexed. Anyway
1057 they got frustrated with this and began working on
1058 thier own project in thier spare time called UNICS :
1059 Uniplexed Information and Computing service.

1060 As time went on people only remebered the acronym
1061 and forgot what the ending CS stood for and it got
1062 transformed into UNIX.

1063 By this time the C programming language was ma-
1064 ture enough that they were able to write this entire
1065 program in just C.

1066 **UNIX vs. LINUX**

1067 This one has been a little muddled in my head for
1068 a while so I thought I would finally write it down here
1069 to explain it to myself.

1070 So linux is an Operating System and a Kernel.

9. Linux Kernel vs Os

1071 **T**his
-§-

-§-

1072 wik (????b)lik (????)lin (????a)wik (????a)lin
1073 (????b)

10. References

- 1074 Likegeeks Filesystem. ????. Accessed: 2019-12-25.
- 1075 Linux Filesystem Explained. ????.a. Accessed: 2019-
1076 12-25.
- 1077 Linuxtopia. ????.b. Accessed: 2019-12-25.
- 1078 Wikipedia : POSIX. ????.a. Accessed: 2019-12-25.
- 1079 Wikipedia Filesystem. ????.b. Accessed: 2019-12-25.